

A Comparison of Distributed Data Processing Systems

Vincent M Chen
Information and Computer science Department
University of Oregon
applekey@cs.uoregon.edu

I. INTRODUCTION

The explosion of data in the past decade has lead to many questions in finding the most efficient method in order to analyze the data being generated. Best put by Google's CEO, Eric Schmidt, "There were 5 exabytes of information created by the entire world between the dawn of civilization and 2003. Now that same amount is created every two days.", more precisely, in 2003, 6.8 exabytes were being created every 2 days,[1]. With such a large volume of throughput, trivial tasks such as locating a users' account or finding all users within the same group can't be found in meaningful time with a single computers nor using sequential algorithms.

In order to solve this problem, a very important insight is to notice that a large majority of the data being created is independent from each other, ex. different users on an on-line shopping platform where every shopper has their own independent profile, shopping cart and payment information. Independence within the data means that parallel computation can be used, where instead of a single computer processing data sequentially, the job is divided into many sub tasks and processed in parallel. This allows the computation of a large data set to be divided among the processing power of multiple computers.

These parallel computation systems have existed since the 1980's, the most common of which are database systems. In the time since, many database systems have evolved into both parallel and distributed varieties. As well, alternative computation frameworks introduced recently such as map reduce or Apache Spark have taken a similarly approach to that used by databases, parallel computation. The goal of this survey paper is to firstly introduce these different data query systems and their variants. Secondly, discuss the common approaches all of them use to speed up computation and finally highlight unique aspects of each system.

II. BASIC DESCRIPTION OF FUNCTIONALITY

A data processing solution should satisfy the following functionality:

1. Guarantee high degree of fault tolerance.
2. Compute a user defined query in a reasonable time.

A. Fault Tolerance

Fault tolerance is the most important guarantee that a computation framework provides. Without consistency, the data

being returned from the query cannot be used for any meaningful analysis. Fault tolerance means recovery from failure from multiple sources

1. Network fault, communication between compute nodes is unreliable and data is lost in between
2. System fault, OS failure
3. Hardware fault, compute computer hard disk fails
4. User fault, user query has code that leaves query engine in a bad state

In order for user defined queries to complete in a reasonable time, two levels of parallelization can be applied.

B. Parallel Computation

The framework can process independent queries in parallel. For example, searching for all students who owe more than 10,000 dollars in student loans can be queried in parallel, each student can be treated independently.

C. Distributed Computation

The framework is able to distribute the data set across multiple compute machines. Each machine might contain independent data, for example, one machine contains all students with first names starting A to J and another from K to Z. Another case is multiple machines hold data that is related to one another. For example in figure 1, machine A might hold all student's GPA and Awards and another machine holds the student's student loan amount. If the query were to find all students that had more than 10,000 in student loans but also a 4.0 GPA in order to give them a financial award, then each query per student would need to look at the data in both compute machines.

III. DATABASES MANAGEMENT SYSTEMS

Databases management systems are frameworks that are designed to store data in tabular tables with a set schema. DBMS systems use relational queries to perform CRUD (Create, read, update and delete) operations on the stored data.

A. Parallel DBMS

Relational queries are ideal suited to run parallel since they each operation declares independent operations on indepen-

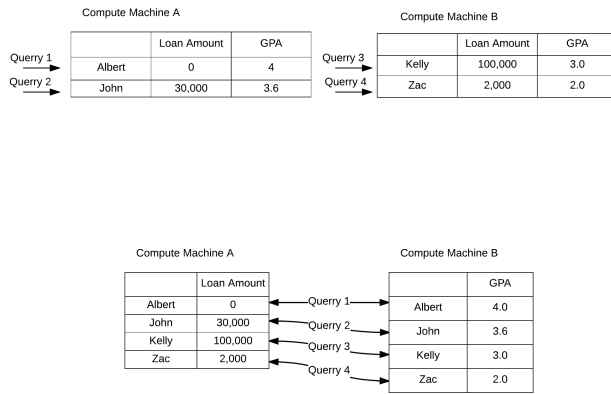


Fig. 1. Distributed Data Dependency

dent data sets [2]. By partitioning the data among multiple processors, a single query can be carried out in parallel and the result merged into a final step. In figure 2, an independent scan (ex, look for key word) is applied and the results are sorted (binned by first letter) and then the result of each independent operation is merged into the final result.




Fig. 2. Distributed Data Dependency

B. Distributed DBMS

Distributed DBMS systems require an additional layer of abstraction on top of Parallel DBMS. Where as in parallel DBMS systems, all queries are run on the same machine sharing the same memory space. Distributed DBMS systems need to consider the additional factor of latency when access data that does not reside in the current compute machine. Additional, distributed DBMS systems need to consider what impact does the memory layout have in preventing easily adding more compute nodes.

Generally speaking, there are three main types of memory architectures for distributed DBMS systems, shared-nothing, shared-memory, shared-disk.

1. Shared-Memory: All compute machines share access to a global memory as well as disk storage.

2. Shared-Disk: Each compute machines have their own memory, but share global disk.

3. Shared Nothing: Each compute machine has their own memory and disk.

C. Fault Tolerance

DBMS systems have very mature and well-studied failure mechanisms. For the purpose of this survey paper, the exact details of each will not be discussed. DBMS systems take advantage of the schema and transactions implement restart at different granularity levels. Through the use of larger compute node level transactions as well as much smaller procedural level transactions, DBMS systems are able to advantage of the work that has already being completed and restart at the most optimal point before the crash.

IV. MAP REDUCE

Map Reduce is a general distributed computation framework introduced by Google in 2004 [3]. Map reduce is different from distributed DBMS systems in that much more general computations can be defined, not limited by the query language of a rigid schema as is the case with DBMS systems. Map Reduce, like DBMS systems have fault tolerance mechanisms but they operate differently than DBMS systems due to the architectural differences.

A. Mapper and Reducer

Map reduce operates in two steps, the map step and then the reduce step. Both steps are independent in that the reduce step will not start until all mappers have completed. In the map stage, the data is partitioned equally among N mappers. Input data arrives to the mapper as a list of (key, value) pairs, the mapper applies a user-defined map operation. The map operation produces an intermediate (key, value) pairs which are the results of the map operations, these results are written and saved to the file system.

The reducer is again a user defined operation that processes the intermediate values from the mapper grouped by key. The reducer can produce zero or an aggregate of all values that map to the same key.

Figure 3 gives an execution diagram.

B. Hadoop

There are two very popular frameworks that implement the map reduce concept, these are the closed source Google Map Reduce and open source version Hadoop Map Reduce. This paper will focus on Hadoop instead of Google Map Reduce

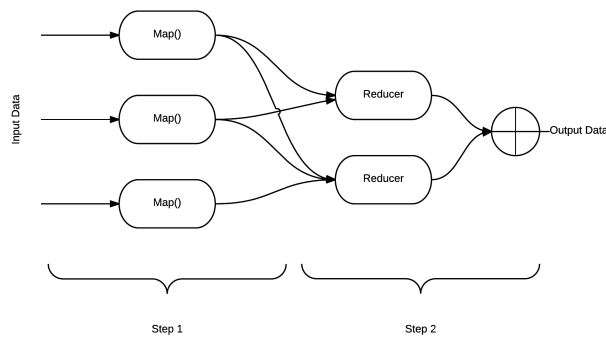


Fig. 3. Map Reduce Program Execution

since the complete details of Google Map Reduce and relatively unknown, however both frameworks function similarly.

Hadoop consist of three layers:

1. Application layer
2. Map Reduce layer
3. HDFS layer

The HDFS layer is Hadoop's storage layer. The HDFS structures data in blocks that is managed by a master node, when the map reduce runs, the master node will assign each individual mapper data chunks as broken down by the HDFS. More importantly, the HDFS layer has the is one major component Hadoop Map Reduce's mechanism for fault tolerance. The application layer sits on top of the map reduce layer. The application layer consists of several additions to the standard map reduce framework, these include high level languages, database adaptations. These application layer additions will be discussed later on in this survey paper.

C. Parallel and Distributed Computation

The main advantage of the map reduce concept is that there are no dependency between mappers or mappers. Map reduce therefore is extremely scalable, adding additional nodes either through adding additional cores or compute nodes is a relatively simple task. This design is very similar to that of the share nothing architecture found in parallel DBMS systems.

D. Fault Tolerance

Hadoop's fault tolerance mechanism is extremely basics, it achieves fault tolerance through replication. When a map reduce job is invoked, the input data is triple replicated on the HDFS. When the mapper or reducer has completed, again the intermediate/final result is triple duplicated. When the master

node, which is responsible for keeping track fo map/reduce jobs notices that a node has failed, the master node will reschedule the map/reduce job with the input/intermediate data.

V. DBMS VS MAP REDUCE

Map reduce and distributed DBMS systems are very similar in that both are parallel data processing frameworks, both support roughly the same feature sets, performance, fault tolerance, user-combustibility. However, map reduce is far more flexible, where as DBMS require data to be organized into schema map reduce can be applied to any variety of data set. This flexibility is not without its trade-offs though, in general is found to be 15 - 20 percent slower than that of map reduce in a variety of general tasks [4]. This section will cover the differences between the two and present a performance study between the two.

A. Map Reduce Advantages

1. Simplicity in Framework

Programming map reduce is easier to use than query languages such as SQL for two reasons. Firstly, there are no complex language constructs as is the case with SQL, and secondly, the programmer does not need to understand a table's schema and data layout.

2. Flexibility

The lack of schema's allow data to be organized in any fashion. The programmer does not need to massage the input data into a rigid format.

3. Storage Agnostic

Since there are no set schema's, data can easily be moved from one storage framework (representation) to another. There is very little needed to re-run the same map/reduce job. For DBMS systems, there would have to be a translation job between two different storage frameworks to ensure that the integrity of the data is kept constant.

B. Map Reduce Disadvantages

1. Lack of High Level Language

While map reduce is extremely to program due to the simplistic nature of the logical map and reduce function, the lack of a high level language means that there is very little in terms of portability when multiple map reduce stages are chained together. Take for example, a very common query in SQL, select * from table where column 1 > value 1 and column 2 > value 2. In map reduce, there would be two stages, the first stage would filter by column 1 and the second column 2. The programmer would need to know to pass data from column 1 in first, then pass in data from column 2. The programmer cannot pass in an entire row with multiple columns because map reduce only supports single (key,value) pairs. In this regards, map reduce has often being compared with assembly

programming, it is difficult to re-use a preexisting map reduce stage without knowing entirely the execution flow for it.

2. Rigid Data Flow

Map reduce runs in two discrete steps, the reduce step does not start until all the mapper steps have completed. Since map reduce is a distributed system running on any kind of network, the possibility that there is a single straggler node is entirely not out of the ordinary. In this situation, all compute nodes would have to wait on the straggler node. DBMS systems do not have this constraint, using the Schema and query logic, the scheduler will schedule as many local reduces as possible from free compute nodes, this results faster execution times [5].

3. Lack of Schema or Indexes

One of map reduce's biggest advantages, flexibility is also one of its largest disadvantages. While not having a schema means that map reduce jobs have very fast initiation times, on the other searching operations are considerably slower. A lack of schema also means that between multiple map reduce stages, data has to be parsed again and again. In addition, the lack of a schema means that compile time optimizations cannot be made, only runtime optimizations as well, reductions in logic between multiple map reduce stages.

C. Difference in Programming Model

Map reduce and DBMS represent an old discussion on two different views in methods to access data [4]. These are:

1. Relational view, state what you want rather than presenting an algorithm; this represents map reduce.
2. Codasyl, an algorithm for data access; this represents query languages such as SQL.

At the time, it was concluded that relational are easier to understand and modify with Codasyl being criticized for being "too assembly similar to assembly".

VI. MAP REDUCE VERSUS DBMS PERFORMANCE

Hadoop Map Reduce and Parallel DBMS systems compared in performance for three category of tasks on a 100 node cluster. The execution times between the two are presented in table I.

A. System Configurations

Both map reduce and DBMS systems were deployed on a 100 node cluster. Each node has a single 2.40 GHz Intel Core 2 Duo with 4GB of ram and two 250GB SATA-I hard disks. All 100 nodes were connected using two gigabyte 128 Gbps Ethernet switch with 50 nodes per switch.

B. Grep Task

The first task was the grep task as was presented in the original Google Map Reduce paper. This task involved searching for a specific keyword that occurred once every 1000 (key, value) pairs.

C. Web Log Task

D. Join Task

	Hadoop	DBMS-x
Grep	284s	194s
Web Log	1,146s	740s
Join	1,1158s	32s

TABLE I
HADOOP VS DBMS-X PERFORMANCE

E.

REFERENCES

- [1] John Gantz and David Reinsel, "The digital universe decade-are you ready," *IDC iView*, 2010.
- [2] David DeWitt and Jim Gray, "Parallel database systems: the future of high performance database systems," *Communications of the ACM*, vol. 35, no. 6, pp. 85–98, 1992.
- [3] Jeffrey Dean and Sanjay Ghemawat, "Mapreduce: Simplified data processing on large clusters, OSDI'04: Sixth symposium on operating system design and implementation, san francisco, ca, december, 2004," *S. Dill, R. Kumar, K. McCurley, S. Rajagopalan, D. Sivakumar, ad A. Tomkins, Self-similarity in the Web, Proc VLDB*, 2001.
- [4] Andrew Pavlo, Erik Paulson, Alexander Rasin, Daniel J Abadi, David J DeWitt, Samuel Madden, and Michael Stonebraker, "A comparison of approaches to large-scale data analysis," in *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data*. ACM, 2009, pp. 165–178.
- [5] Hung-chih Yang, Ali Dasdan, Ruey-Lung Hsiao, and D Stott Parker, "Map-reduce-merge: simplified relational data processing on large clusters," in *Proceedings of the 2007 ACM SIGMOD international conference on Management of data*. ACM, 2007, pp. 1029–1040.