

# A Comparison of Distributed Data Processing Systems

Vincent M Chen  
Information and Computer science Department  
University of Oregon  
*applekey@cs.uoregon.edu*

## I. INTRODUCTION

The explosion of data in the past decade has lead to many questions in finding the most efficient method in order to analyze the data being generated. Best put by Google's CEO, Eric Schmidt, "There were 5 exabytes of information created by the entire world between the dawn of civilization and 2003. Now that same amount is created every two days.", more precisely, in 2003, 6.8 exabytes were being created every 2 days,[1]. With such a large volume of throughput, trivial tasks such as locating a users' account or finding all users within the same group can't be found in meaningful time with a single computers nor using sequential algorithms.

In order to solve this problem, a very important insight is to notice that a large majority of the data being created is independent from each other, ex. different users on an on-line shopping platform where every shopper has their own independent profile, shopping cart and payment information. Independence within the data means that parallel computation can be used, where instead of a single computer processing data sequentially, the job is divided into many sub tasks and processed in parallel. This allows the computation of a large data set to be divided among the processing power of multiple computers.

These parallel computation systems have existed since the 1980's, the most common of which are database systems. In the time since, many database systems have evolved into both parallel and distributed varieties. As well, alternative computation frameworks introduced recently such as map reduce or Apache Spark have taken a similarly approach to that used by databases, parallel computation. The goal of this survey paper is to firstly introduce these different data query systems and their variants. Secondly, discuss the common approaches all of them use to speed up computation and finally highlight unique aspects of each system.

## II. BASIC DESCRIPTION OF FUNCTIONALITY

A data processing solution should satisfy the following functionality:

1. Guarantee high degree of fault tolerance.
2. Compute a user defined query in a reasonable time.

### A. Fault Tolerance

Fault tolerance is the most important guarantee that a computation framework provides. Without consistency, the data

being returned from the query cannot be used for any meaningful analysis. Fault tolerance means recovery from failure from multiple sources

1. Network fault, communication between compute nodes is unreliable and data is lost in between
2. System fault, OS failure
3. Hardware fault, compute computer hard disk fails
4. User fault, user query has code that leaves query engine in a bad state

In order for user defined queries to complete in a reasonable time, two levels of parallelization can be applied.

### B. Parallel Computation

The framework can process independent queries in parallel. For example, searching for all students who owe more than 10,000 dollars in student loans can be queried in parallel, each student can be treated independently.

### C. Distributed Computation

The framework is able to distribute the data set across multiple compute machines. Each machine might contain independent data, for example, one machine contains all students with first names starting A to J and another from K to Z. Another case is multiple machines hold data that is related to one another. For example in figure 1, machine A might hold all student's GPA and Awards and another machine holds the student's student loan amount. If the query were to find all students that had more than 10,000 in student loans but also a 4.0 GPA in order to give them a financial award, then each query per student would need to look at the data in both compute machines.

## III. DATABASES MANAGEMENT SYSTEMS

Databases management systems are frameworks that are designed to store data in tabular tables with a set schema. DBMS systems use relational queries to perform CRUD (Create, read, update and delete) operations on the stored data.

### A. Parallel DBMS

Relational queries are ideal suited to run parallel since they each operation declares independent operations on independent data sets [2]. By partitioning the data among multiple processors, a single query can be carried out in parallel and the result merged into a final step. In figure 2, an independent scan (ex, look for key word) is applied and the results

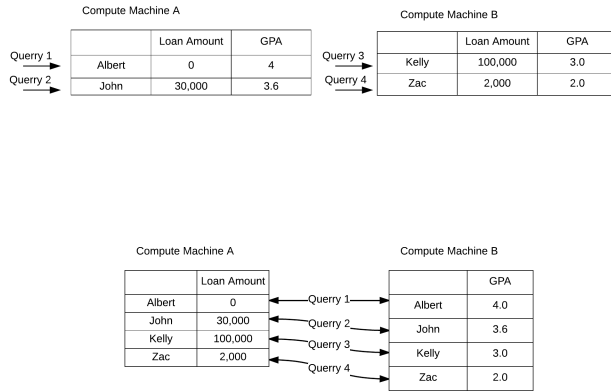


Fig. 1. Distributed Data Dependency

are sorted (binned by first letter) and then the result of each independent operation is merged into the final result.




Fig. 2. Distributed Data Dependency

### B. Distributed DBMS

Distributed DBMS systems require an additional layer of abstraction on top of Parallel DBMS. Whereas in parallel DBMS systems, all queries are run on the same machine sharing the same memory space. Distributed DBMS systems need to consider the additional factor of latency when access data that does not reside in the current compute machine. Additionally, distributed DBMS systems need to consider what impact does the memory layout have in preventing easily adding more compute nodes.

Generally speaking, there are three main types of memory architectures for distributed DBMS systems, shared-nothing, shared-memory, shared-disk.

1. Shared-Memory: All compute machines share access to a global memory as well as disk storage.
2. Shared-Disk: Each compute machines have their own memory, but share global disk.
3. Shared Nothing: Each compute machine has their own memory and disk.

### C. Fault Tolerance

## IV. MAP REDUCE

Map Reduce is a general distributed computation framework introduced by Google in 2004 [3]. Map reduce is different from distributed DBMS systems in that much more general computations can be defined, not limited by the query language of a rigid schema as is the case with DBMS systems. Map Reduce, like DBMS systems have fault tolerance mechanisms but they operate differently than DBMS systems due to the architectural differences.

### A. Mapper and Reducer

Map reduce operates in two steps, the map step and then the reduce step. Both steps are independent in that the reduce step will not start until all mappers have completed. In the map stage, the data is partitioned equally among N mappers. Input data arrives to the mapper as a list of (key, value) pairs, the mapper applies a user-defined map operation. The map operation produces intermediate (key, value) pairs which are the results of the map operations.

The reducer is again a user defined operation that processes the intermediate values from the mapper grouped by key. The reducer can produce zero or an aggregate of all values that map to the same key.

Figure 3 gives an execution diagram.

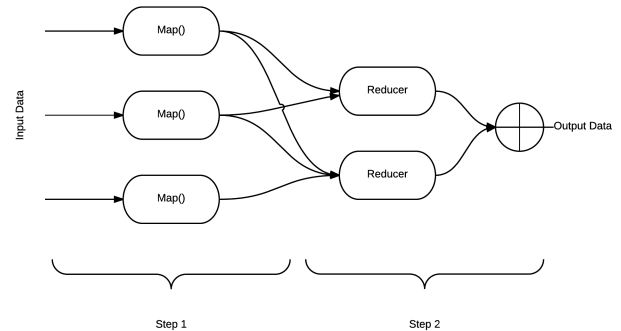


Fig. 3. Map Reduce Program Execution

### B. Hadoop

There are two very popular frameworks that implement the map reduce concept, these are the closed source Google Map Reduce and open source version Hadoop Map Reduce. This paper will focus on Hadoop instead of Google Map Reduce since the complete details of Google Map Reduce and relatively unknown, however both frameworks function similarly.

Hadoop consist of three layers:

1. Application layer
2. Map Reduce layer

### 3. HDFS layer

#### *C. Fault Tolerance*

Analogous to parallel DBMS, the map and reduce step are very similar to a filter and reduce.

#### REFERENCES

- [1] John Gantz and David Reinsel, "The digital universe decade-are you ready," *IDC iView*, 2010.
- [2] David DeWitt and Jim Gray, "Parallel database systems: the future of high performance database systems," *Communications of the ACM*, vol. 35, no. 6, pp. 85–98, 1992.
- [3] Jeffrey Dean and Sanjay Ghemawat, "Mapreduce: Simplified data processing on large clusters," *OSDI'04: Sixth symposium on operating system design and implementation*, san francisco, ca, december, 2004," *S. Dill, R. Kumar, K. McCurley, S. Rajagopalan, D. Sivakumar, and A. Tomkins, Self-similarity in the Web, Proc VLDB*, 2001.